

# SARSA : Trace accumulation vs Trace Replacement

## Algorithm

```
Initialize Q(s,a) arbitrarily and e(s,a)=0, for all s,a
Repeat(for each episode)
  Initialize s,a
  Repeat(for each step of episode)
    Take action a, observe r, s'
    Choose a' from s' using the behavior policy
     $\delta = r + \gamma Q(s',a') - Q(s,a)$ 
    e(s,a) = 1
    For all s,a:
       $Q(s,a) = Q(s,a) + \alpha \delta e(s,a)$ 
      e(s,a) =  $\gamma \lambda e(s,a)$ 
    s = s'
    a = a'
  until s is terminal
```

## Trace Accumulation

The entire SARSA learning algorithm remains same, but the trace updation step in this case is :  $\text{trace}[s, a] = \text{trace}[s, a] + 1$

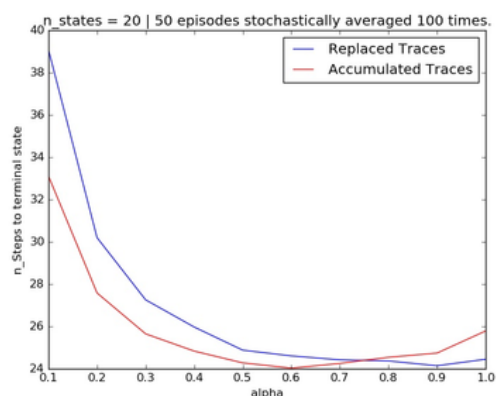

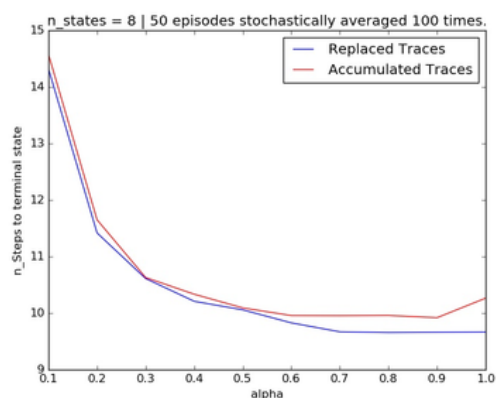

## Trace Replacements

The entire SARSA learning algorithm remains same, but the trace updation step in this case is :  $\text{trace}[s, a] = 1$

## Comparison

For the given MDP, since we are given the reward only on reaching the terminal state, **accumulating traces** is detrimental to learning as individual step rewards are unknown. **Replacing traces** works well with the **highly stochastic** nature of the MDP and learns in a **prior-free** (unbiased from previous decisions) fashion and thus converges to termination **faster** than the former. We can empirically see using the following figures which clear show that replacing traces **outperforms** accumulating traces for relatively high alphas.

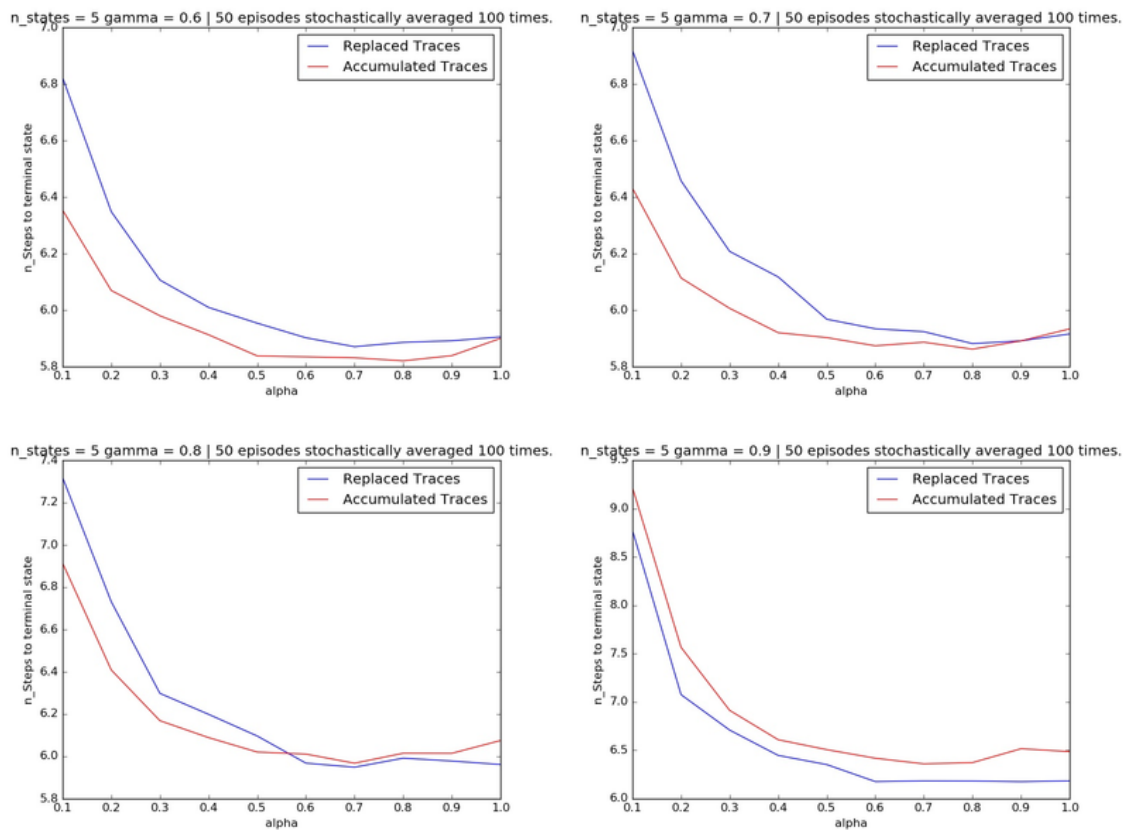
Following plots with all parameters except the  $n\_states$  , we can clearly see that **replacing traces > accumulating traces** for relatively high alphas.



## Conclusion

For a given  $\gamma$ ,  $\alpha$  and  $\epsilon$  replacing traces outperforms accumulating traces. However we can empirically see this advantage disappears as we increase the number of states / decrease the value of  $\gamma$ . The intuition behind this can be seen by the fact that with increment in the number of states / decrease in  $\gamma$ , the reward-delay increases and the MDP becomes more non-Markovian. Since eligibility traces (replacing) are used as a trade-off between TD and MC methods to overcome reward-delays, we see that it **loses its advantages for very long and less discounted** MDPs as the one used in this problem.

In all these plots we increase the value of  $\gamma$  from 0.5 to 0.9 in steps of 0.1 :



## References

- [http://www-all.cs.umass.edu/pubs/1995\\_96/singh\\_s\\_ML96.pdf](http://www-all.cs.umass.edu/pubs/1995_96/singh_s_ML96.pdf)
- [https://www.tu-chemnitz.de/informatik/KI/scripts/ws0910/ml09\\_7.pdf](https://www.tu-chemnitz.de/informatik/KI/scripts/ws0910/ml09_7.pdf) (Best explanation for accumulation traces)
- <http://www.karanmg.net/Computers/reinforcementLearning/finalProject/KaranComparisonOfSarsaWatkins.pdf>

## Running code

As suggested in the instructions my code can be run using the files in the `src` directory:

- `python sarsa.py <n_states> <gamma>`
- `./sarsa.sh` to run all the experiments

Plots (similar to the ones I have provided in `imgs` directory) will be generated

## Additional Notes

- Due to stochasticity of the experiments, I chose to average the results for first  $n_{episodes} = 50$  over a large number ( $n_{trials} = 100$ ) iterations to retain the credibility of the results obtained.
- I have chosen 10  $\alpha$  values ranging from 0.1 to 1 (both inclusive) with the  $\alpha[i] = 0.1 * i$ .
- The values for  $\gamma$  and  $\epsilon$  were empirically chosen so as to run the experiment quickly.